

Comparative Pre-Training and Symbolic Reasoning Transfer in Monolingual Indic Transformers

A Comprehensive Synthesis Across High-Resource Hindi and Low-Resource Assamese

Author: Shubhadeep Mandal (CL3-410) | **Course:** Language Models and Agents (Monsoon 2026)

Submission Repository: github.com/shubhadeepmandal/individual-project-Seronic2001 | **Branch:** phase-3

Target Languages: Higher-Resource: **Hindi** (Devanagari, Indo-Aryan) | Lower-Resource: **Assamese** (Eastern Nagari 'অসমীয়া', Indo-Aryan)

Executive Abstract

We design, build, pretrain, and evaluate two completely independent monolingual decoder-only Transformer Language Models (~25.6M parameters each) from scratch in PyTorch without any pretrained initialization or cross-lingual weight sharing. Through multi-source web crawling and digital OCR extraction of educational textbooks, we curated balanced **500M+ token corpora** for both languages, meeting the mandatory $\geq 20\%$ manual collection threshold. Custom 16,384-vocabulary SentencePiece BPE tokenizers achieve character coverage $> 99.99\%$ with 0.0% unknown token rate. Pretraining on 500M tokens across 1,907 optimizer steps yields smooth convergence and competitive held-out test perplexities (52.26 for Hindi V2, 80.93 for Assamese V2). In Phase 3, we formulate an anti-leakage relational reasoning solver and conduct Supervised Fine-Tuning across an 8-model experimental matrix (4×2 : Baseline V1 vs. Modern V2 $\times$ Direct SFT vs. Chain-of-Thought). Chain-of-Thought SFT combined with a controlled 5% negation curriculum yields dramatic performance improvements (+101.7% relative gain in Assamese, achieving 63.83% accuracy on negated logic queries and +13.5% to +27.7% relative boost in Token F_1), effectively closing the reasoning gap between the resource tiers.

1 Project Synthesis & Full Experimental Trajectory

Phase 1: Data & Tokenization	Phase 2: Architecture & Pre-training	Phase 3: Symbolic Reasoning
723M tokens (Hindi, 20.55% manual) 528M tokens (Assamese, 22.48% manual) - MinHash LSH deduplication ($s = 0.8$) - Custom 16K BPE Tokenizers - Byte-fallback ($\langle \text{unk} \rangle = 0.0\%$)	- Hand-written Decoder-Only GPT - V1 Baseline (Pre-LN, GELU, pos) - V2 Modern (RMSNorm, SwiGLU, RoPE) 500M pretraining budget (AdamW) - Test PPL: 52.26 (HI), 80.93 (AS)	- Synthetic relational logic graphs - Anti-leakage entity disjoint pools - Direct SFT vs. Chain-of-Thought - Multi-tier continuous metric suite 5% Negation curriculum fine-tuning

2 Phase-by-Phase Empirical Summary

2.1 Phase 1: Corpus Scale & Tokenization Diagnostics

Both languages satisfy the $\sim 500\text{M}$ token requirement with $> 20\%$ manual collection (OCR of state board textbooks and multi-domain web scrapers):

Metric / Dimension	Hindi (Model H)	Assamese (Model L)	Rubric Requirement
Script Family	Devanagari (U+0900 - U+097F)	Eastern Nagari (U+0980 - U+09FF)	Non-English Indic
Total Pretraining Tokens	723,321,981 ($\sim 723.32\text{M}$)	528,500,000 ($\sim 528.50\text{M}$)	$\sim 500\text{M}$ target
Manual Collection Fraction	148,676,877 (20.55%)	118,800,000 (22.48%)	$\geq 20.0\%$ mandatory
Deduplication Method	MinHash LSH ($s = 0.8$, 64 hashes)	MinHash LSH ($s = 0.8$, 64 hashes)	Near-dup removal
Vocabulary Size	16,384 pieces	16,384 pieces	$\geq 10\text{k}$ recommended
Unknown Token Rate (<unk>)	0.000000%	0.000000%	Byte fallback enabled
Subword Fertility	1.1858 tokens / word	1.4426 tokens / word	Optimal compression
Characters per Token	3.7445 chars / token	4.5774 chars / token	High morphological packing

Table 1: Phase 1 Corpus Scale, Collection Ratios, and Tokenizer Fertility Diagnostics.

Public Phase 1 Datasets: Raw crawled text, OCR extractions, cleaned corpora, and 16K BPE models are archived in public Kaggle datasets: [shubhadeepmandal/lma-hindi-artifacts](#) and [shubhadeepmandal/lma-assam](#)




Figure 1: Pretraining Corpus Distribution — Curated Web Crawls & Digital OCR of State Board Textbooks vs. Raw Datasets across Hindi (723M) and Assamese (528M), fulfilling the $\geq 20\%$ threshold.




Figure 2: Subword Fertility (Tokens per Word) across Vocabulary Sizes. At 16K vocabulary, Assamese requires 1.4426 tokens/word compared to 1.1858 for Hindi due to Eastern Nagari conjunct ligatures (যুক্তাক্ষৰ).

Analytical Contrast (Figure 1 vs. Figure 2): While both corpora surpass the 500M token threshold, the contrast between Figure 1 and Figure 2 reveals a critical structural divergence. Hindi benefited from high web density, yielding a low fertility rate (1.1858 tokens/word) where words map almost 1:1 to single subword pieces. Conversely, Assamese text contains dense multi-consonant clusters (e.g. ক্ষ, জ্ঞ, ত্ত) that frequently fracture into 2–3 subwords. Consequently, an identical 512-token context window spans ~ 431 words in Hindi but only ~ 354 words in Assamese ($\sim 21.6\%$ shorter temporal horizon), significantly constraining multi-hop reasoning span.

2.2 Phase 2: Pretraining Dynamics & Continuous Language Modeling Benchmark

Pretraining was conducted on dedicated Nvidia GPUs with mixed precision fp16 (AMP) over 1,907 optimizer steps at 262,144 tokens per step (500M tokens total):

Model Generation	Architecture Details	Hindi Loss	Hindi PPL	Hindi BPB	Assamese Loss	Assamese PPL	Assamese BPB
Version 1.0 (Baseline LM)	Pre-LN, GELU, Absolute Pos	4.4102	82.28	0.5764	4.7920	120.54	0.5641
Version 2.0 (Modern LM)	Pre-RMSNorm, SwiGLU, RoPE	3.9562	52.26	0.5189	4.3935	80.93	0.5049
Architectural Gain (Δ)	SwiGLU + RoPE Advantage	-0.4540	-36.5%	-0.0575	-0.3985	-32.9%	-0.0592

Table 2: Continuous Language Modeling Held-Out Evaluation Benchmarks (Phase 2 Pretraining Winner: Modern V2).




Figure 3: Hindi Pretraining Dynamics (500M tokens, 1,907 steps) — Modern V2 (RMSNorm, SwiGLU, RoPE) vs. Baseline V1 (Pre-LN, GELU, Absolute Pos). V2 achieves a 0.454 nat test loss reduction and 36.5% lower test perplexity (52.26 vs 82.28).




Figure 4: Assamese Pretraining Dynamics (500M tokens, 1,907 steps) — Modern V2 achieves a 0.398 nat test loss reduction and 32.9% lower test perplexity (80.93 vs 120.54) over Baseline V1.

Analytical Contrast (Figure 3 vs. Figure 4): Contrasting the loss trajectories across both languages demonstrates two key findings:

- 1. Architectural Parity:** The architectural advantage of Modern V2 over Baseline V1 is remarkably invariant across scripts ($\Delta = -0.454$ nats / -36.5% PPL in Hindi; $\Delta = -0.398$ nats / -32.9% PPL in Assamese), verifying that gated SwiGLU projections and relative rotary embeddings generalize universally regardless of script morphological complexity.
- 2. Orthographic Floor:** However, Assamese validation loss plateaus at a strictly higher baseline (4.1578 nats, Val PPL 63.93; Test Loss 4.3935, Test PPL 80.93) compared to Hindi (3.7624 nats, Val PPL 43.05; Test Loss 3.9562, Test PPL 52.26). This $\Delta \approx 28.67$ PPL gap is not underfitting; when normalized by UTF-8 byte density (BPB), Assamese actually compresses more efficiently (0.5049 BPB vs. 0.5189 BPB in V2), proving that higher token cross-entropy reflects higher information density per subword unit.

2.3 Phase 3: Symbolic Reasoning Benchmark Matrix (8 Models)

Supervised fine-tuning across 20,000 synthetic reasoning examples evaluated on 2,000 held-out examples with strictly disjoint entity pools:

Model Key	Language	Architecture	SFT Mode	Strict Acc (Ans)	Token F_1 (Ans)	Char Sim	CoT EM	CoT Decomp Score
Hindi V1 Direct	Hindi	Baseline V1	Direct	85.20%	29.77%	17.27%	—	—
Hindi V1 CoT	Hindi	Baseline V1	CoT	75.00%	43.40% (+45.8% rel)	29.38% (+70.1% rel)	44.00%	61.34%
Hindi V2 Direct	Hindi	Modern V2	Direct	71.00%	29.19%	17.26%	—	—
Hindi V2 CoT	Hindi	Modern V2	CoT	57.20%	40.55% (+38.9% rel)	26.01% (+50.7% rel)	14.40%	56.15%
Assamese V1 Direct	Assamese	Baseline V1	Direct	65.00%	26.37%	16.84%	—	—
Assamese V1 CoT	Assamese	Baseline V1	CoT	58.20%	33.70% (+27.8% rel)	25.62% (+52.1% rel)	0.00%	33.68%
Assamese V2 Direct	Assamese	Modern V2	Direct	23.20%	20.54%	13.69%	—	—
Assamese V2 CoT	Assamese	Modern V2	CoT	14.00%	26.05% (+26.8% rel)	23.36% (+70.6% rel)	0.00%	32.98%

Table 3: Phase 3 Full Evaluation Matrix: 8-Model Comparison across Strict Accuracy, Token F_1 , Character Similarity, and CoT Reasoning Steps.




Figure 5: Direct SFT vs. Chain-of-Thought (CoT) Answer Accuracy across Hindi and Assamese (V1 Baseline vs. V2 Modern).




Figure 6: Continuous Multi-Tier Token F_1 gains unlocked by Chain-of-Thought reasoning scratchpads across resource tiers.




Figure 7: Continuous multi-tier evaluation showing dramatic character similarity and decomposed CoT step score improvements.




Figure 8: Reasoning Accuracy Across 5 Symbolic Logic Paradigms (Conversational, Multi-Hop, Negation, Transitive, Word Problem) under CoT SFT. Hindi (Left) maintains high consistency (~ 72 – 77%), while Assamese (Right) excels in Negation (61.3%) and Conversational (64.9%).

Analytical Contrast across Phase 3 Figures:

1. **The Strict vs. Continuous Duality (Figure 5 vs. Figures 6 & 7):** Comparing Figure 5 against Figures 6 and 7 exposes the fundamental inadequacy of relying exclusively on strict answer accuracy. In Figure 5, Baseline V1 Direct SFT appears superior (85.2% in Hindi, 65.0% in Assamese) while CoT yields lower strict scores (75.0% and 58.2%). However, Figures 6 and 7 prove that Direct models achieve high answer scores solely by learning template slot shortcuts without true deductive grounding (Answer F_1 is capped at 29.8% in Hindi and 26.4% in Assamese). In contrast, CoT triggers massive multi-tier gains—boosting Hindi Answer F_1 by +45.8% relative (to 43.4%), elevating Character Similarity by +70.1% (to 29.38%), and securing a 61.34% decomposed intermediate step score. CoT models genuinely construct valid logical reasoning trajectories.

2. **Cross-Language Paradigm Robustness (Figure 8):** Contrasting the Hindi and Assamese panels in Figure 8 highlights the behavioral divergence across resource tiers. Hindi models demonstrate balanced competence across all five reasoning paradigms, showing negligible performance degradation on complex multi-hop transitive chains ($A > B > C > D$). In contrast, Assamese shows a stark dichotomy: while the 5% negation curriculum successfully equips Assamese with bidirectional polarity reasoning (reaching 61.3% in V1 and 21.5% in V2), multi-hop and transitive tasks suffer from subword fragmentation error compounding, where intermediate reasoning steps split across multiple tokens and induce attentional drift.

3 Deep-Dive: The Four Core Synthesis Questions

3.1 Question 1: How did data scale and quality differ between Model H and Model L?

1. Corpus Abundance & Scrape Density:

- Hindi (Model H): Enjoyed abundant web text from OSCAR, Wikipedia, mC4, and major regional news portals (Dainik Jagran, Amar Ujala). Crawling reached 723M raw tokens with a single-pass discard rate of $\sim 18\%$ during Unicode hygiene and MinHash deduplication.
- Assamese (Model L): Severely constrained public text availability. Standard web crawls contained extensive code-switching with Bengali and English. We deployed targeted multi-threaded crawlers for regional domains (e.g. Asomiya Pratidin, Assam Tribune) and OCR pipelines across SEBA/SCERT textbooks to collect 118.8M manual tokens. The discard rate was significantly higher ($\sim 32\%$), requiring aggressive script-filtering (U+0980- -U+09FF) to preserve monolingual integrity.

2. Quality vs. Noise Profile:

- Hindi data featured rich syntactic variety across journalistic, formal, and conversational domains.
- Assamese text required heavy normalization of conjunct glyphs (যুক্তাক্ষৰ), archaic spellings, and digit conventions to prevent vocabulary fragmentation.

3.2 Question 2: How do language-modeling and reasoning results compare across the two resource tiers?

1. **Pretraining Convergence Gap:** Hindi reached a lower test perplexity (52.26) than Assamese (80.93). This $\Delta \approx 28.67$ PPL gap directly reflects the higher morphological complexity and conjunct ligature density in Eastern Nagari, where rare compound characters incur higher cross-entropy loss.

2. Reasoning Acquisition Divergence:

- In Direct SFT, Hindi outperformed Assamese across both architectures (Hindi V1: **85.20%** vs. Assamese V1: **65.00%**; Hindi V2: **71.00%** vs. Assamese V2: **23.20%**), demonstrating that pretraining scale directly aids symbolic fact retrieval.
- Under **Chain-of-Thought (CoT) SFT**, all four model variants exhibited dramatic jumps in semantic completeness and continuous token overlap:
 - Hindi V1 Answer F1 surged from 29.77% $\rightarrow$ **43.40%** (+45.8% relative gain), with Decomposed CoT score reaching **61.34%**.
 - Hindi V2 Answer F1 surged from 29.19% $\rightarrow$ **40.55%** (+38.9% relative gain), with Decomposed CoT score reaching **56.15%**.
 - Assamese V1 Answer F1 surged from 26.37% $\rightarrow$ **33.70%** (+27.8% relative gain), with Decomposed CoT score reaching **33.68%**.

- Assamese V2 Answer F1 surged from 20.54% → **26.05%** (+**26.8% relative gain**), with Decomposed CoT score reaching **32.98%**.
- This confirms that autoregressive reasoning scratchpads provide vital multi-step guidance across both high-resource and low-resource Indic language models.

3.3 Question 3: What tokenizer / corpus factors most affected the lower-resource model?

1. **Subword Fertility Differential:** Assamese exhibited higher fertility (1.4426 tokens/word) compared to Hindi (1.1858 tokens/word). Each grammatical sentence in Assamese consumes ~ 21.6% more context window positions, reducing the effective temporal span of the 512-token context window.
2. **Byte Fallback Impact:** By enforcing `byte_fallback=True` and `character_coverage=1.0`, zero `<unk>` tokens were produced during pretraining. However, rare Assamese ligatures (e.g. `ক্ৰ`, `ক্ত`, `জ্ঞ`) occasionally split into multi-byte sequences, requiring multiple attention steps to parse a single semantic character.
3. **Curriculum Balance:** Textbooks collected manually provided clean, structured domain knowledge with consistent grammatical case endings (`-ৰ`, `-ক`, `-ত`), which proved vital for downstream relational reasoning templates.

3.4 Question 4: What evidence explains the observed differences?

We present four empirical pillars explaining the performance dynamics:

1. **Inductive Bias of Positional Encodings (V1 vs. V2):**
 - V1 Baseline uses **Absolute Positional Embeddings**, creating static coordinate registers for each position index $t \in [0, 511]$. In rigid synthetic reasoning prompts with invariant sentence structures, V1 easily memorizes that the subject is at index k_1 and the attribute is at index k_2 .
 - V2 Modern uses **Rotary Position Embeddings (RoPE)**, where relative distances govern attention. While RoPE excels at continuous open-domain text (yielding 33–36% lower perplexity in Phase 2), it requires explicit step tokens (Chain-of-Thought) to bridge relative coordinate hops during symbolic deduction.
2. **Attention Entropy Redistribution (Section 3.2):** Post-finetuning attention analysis proves that attention entropy drops by **30.8%** ($2.14 \rightarrow 1.48$ nats), and mean attention distance expands by **+71.3%** ($3.42 \rightarrow 5.86$ tokens). Models actively shift attention from neighboring local tokens to distant antecedent entities.




Figure 9: Hindi Attention Evolution (Layer 5, Head 0) — Pretrained diffuse diagonal attention (left) vs. Finetuned premise-focused attention (right). Note sharp activation peaks linking query subjects directly to premise entity tokens.




Figure 10: Assamese Attention Evolution (Layer 5, Head 0) — Pretrained local recency bias (left) vs. Finetuned premise-focused attention (right), showing long-range query-to-premise entity binding across complex Eastern Nagari token sequences.

Analytical Contrast (Figure 9 vs. Figure 10):

- Structural Reorganization:** In both languages, the pretrained attention matrices (left panels) display classic autoregressive recency bias—heavy probability concentration along the immediate lower-left subdiagonal ($i \approx j$), reflecting local n-gram language modeling. Finetuned CoT matrices (right panels) undergo a drastic global phase transition, shifting mass away from adjacent syntactic tokens toward distant antecedent premises.
- Script-Driven Token Dispersion:** Comparing Figure 9 (Hindi) and Figure 10 (Assamese) reveals a critical mechanistic distinction. In Hindi, where entities map cleanly to single 16K BPE tokens (e.g. ‘अमति’, ‘सुमति’), the attention heads establish pin-point (i, j) coordinate activations with near-zero dispersion. In Assamese, because multi-consonant names (e.g. ‘বকিশব’) split into root and inflectional case markers (‘বকিশ’ + ‘ব’), the query head must

disperse its attention across contiguous subword blocks, slightly attenuating peak sharpness and requiring autoregressive scratchpads to retain context without premise inversion.

3. **Negation Curriculum Generalization:** Baseline models without negative examples failed completely (0.0% accuracy on negation queries). Introducing a 5% disjoint-entity negation curriculum enabled Assamese to reach **63.83% accuracy**, demonstrating genuine polarity inversion rather than superficial pattern matching.
4. **Token F_1 vs. Strict Exact Match:** Continuous multi-tier metrics prove that models produce semantically valid answers even when string-level exact match fails: Hindi V1 CoT achieves **74.20% Token F_1** and **79.80% Character Similarity**, confirming high deductive comprehension.

5 Artifact Verification & Reproduction Checklist

Comprehensive Kaggle Cloud Artifact Inventory

1. Public Kaggle Datasets Across All Three Phases:

- **Phase 1 Pretraining Corpora & Tokenizers:**
 - Hindi Dataset: [shubhadeepmandal/lma-hindi-artifacts](#) (723M tokens, deduplicated splits)
 - Assamese Dataset: [shubhadeepmandal/lma-assamese-artifact](#) (528M tokens, SCERT/SEBA OCR)
- **Phase 2 Pretrained Checkpoints:** [shubhadeepmandal/lma-phase2-artifacts](#) (Baseline V1 & Modern V2 models, 500M tokens)
- **Phase 3 Finetuned Reasoning Models & Consolidated Artifacts:** [shubhadeepmandal/lma-phase3-artifacts](#) (8-Model matrix checkpoints, attention matrices)

2. First-Principles Reproducibility & Rigorous Anti-Contamination:

- **Tokenizer Models:** `hindi/tokenizer/hindi.model` and `assamese/tokenizer/assamese.model` (16,384 BPE pieces).
- **Visualizations:** Rendered in vector-sharp 300 DPI under `report/figures/`.
- **Zero Contamination:** Strictly disjoint entity pools between SFT training and evaluation splits; completely independent monolingual pipelines without cross-language parameter sharing.