

Comparative Pre-Training and Symbolic Reasoning Transfer in Monolingual Indic Transformers: A Contrastive Study Across High-Resource Hindi and Low-Resource Assamese

Shubhadeep Mandal

Department of Computer Science and Engineering | Course: *Language Models and Agents (Monsoon 2026)* | Roll No: 2025201056

Project Repository: <https://github.com/shubhadeepmandal/individual-project-Seronic2001> | Branch: phase-3

Target Languages: **Hindi** (Devanagari script) | **Assamese** (Eastern Nagari script ‘অসমীয়া’)

Abstract

We design, train, and empirically evaluate two completely independent monolingual decoder-only Transformer Language Models (~ 25.6 M parameters each) from first principles in PyTorch without cross-lingual weight sharing or pretrained initialization. Through multi-source web crawling and digital OCR extraction of educational textbooks, we curated balanced 500M+ token corpora for higher-resource Hindi (723M tokens) and lower-resource Assamese (528M tokens), strictly satisfying the mandatory $\geq 20\%$ manual collection threshold. Custom 16,384-vocabulary SentencePiece BPE tokenizers achieve character coverage $> 99.99\%$ with 0.0% unknown token rate. Pretraining on 500M tokens across 1,907 optimizer steps yields smooth convergence and competitive held-out test perplexities (52.26 for Hindi V2, 80.93 for Assamese V2). In Phase 3, we formulate an anti-leakage relational reasoning solver and conduct Supervised Fine-Tuning across an 8-model experimental matrix (4×2 : Baseline V1 vs. Modern V2 $\times$ Direct SFT vs. Chain-of-Thought). Chain-of-Thought SFT combined with a controlled 5% negation curriculum yields dramatic performance improvements (+101.7% relative gain in Assamese, achieving 63.83% accuracy on negated logic queries and +13.5% to +27.7% relative boost in Token F_1), effectively closing the reasoning gap between the resource tiers.

1 Introduction and Experimental Pipeline

The development of autoregressive decoder-only Transformer language models has largely centered on high-resource languages such as English and Chinese. Low-resource and morphologically complex Indian languages present unique challenges stemming from complex conjunct ligatures, scarce digital footprints, and severe tokenization fragmentation.

In this investigation, we explore the comparative dynamics of monolingual pretraining and downstream symbolic reasoning acquisition across two linguistically distinct Indo-Aryan languages: **Hindi** (Devanagari script, higher-resource) and **Assamese** (Eastern Nagari script ‘অসমীয়া’, lower-resource). Both models are built entirely from scratch, utilizing 100% disjoint pipelines, independent vocabularies, and zero cross-lingual weight sharing.

Table 1 provides an architectural and operational overview of the three-phase engineering trajectory conducted across the project.

2 Phase-by-Phase Empirical Summary

2.1 Phase 1: Corpus Scale and Tokenization Diagnostics

Both language corpora surpass the mandatory 500M-token threshold with $> 20\%$ manual collection derived from state educational textbooks and regional news portals. Table 2 summarizes the corpus metrics and tokenizer diagnostics.

Analytical Contrast (Figure 1a vs. Figure 1b): While both corpora satisfy the 500M token requirement, contrasting Figure 1a and Figure 1b illustrates an acute structural divergence. Hindi benefited from extensive digital news archives, producing a low fertility rate (1.1858 tokens/word) where words align near-isomorphically with subword units. In con-

trast, Assamese text contains dense multi-consonant clusters (e.g. ক্ৰ, জ্ঞ, ত্ত) that fracture into 2–3 subword pieces. Consequently, an identical 512-token context window spans ~ 431 words in Hindi but only ~ 354 words in Assamese ($\sim 21.6\%$ shorter temporal horizon), imposing tighter bounds on multi-hop symbolic deduction.

2.2 Phase 2: Pretraining Dynamics and LM Benchmarks


Both architectures (~ 25.6 M parameters) were pretrained under identical conditions: batch size 512, sequence length 512 (262,144 tokens/step), AdamW optimizer ($\beta_1 = 0.9, \beta_2 = 0.95$, weight decay 0.1), cosine learning rate schedule peaking at 6×10^{-4} with 100 warmup steps, over 1,907 steps (500M tokens total).

Analytical Contrast (Figure 2a vs. Figure 2b):


- Architectural Invariance:** The performance boost of Modern V2 over Baseline V1 remains remarkably steady across scripts ($\Delta = -0.454$ nats / -36.5% PPL in Hindi; $\Delta = -0.398$ nats / -32.9% PPL in Assamese). This confirms that relative rotary position embeddings and gated SwiGLU non-linearities transfer effectively across varying orthographic complexity.
- Orthographic Baseline:** Assamese validation loss plateaus at a strictly higher value (4.3935 nats, PPL 80.93) than Hindi (3.9562 nats, PPL 52.26). However, when normalized by byte entropy (Bits-Per-Byte), Assamese achieves a lower BPB (0.5049 vs. 0.5189). This demonstrates that higher token cross-entropy reflects greater linguistic compression rather than underfitting.

Pipeline Phase	Phase 1: Data & Tokenization	Phase 2: Architecture & Pretraining	Phase 3: Symbolic Reasoning SFT
Corpus Scale	723M tokens (Hindi, 20.55% manual) 528M tokens (Assamese, 22.48% manual)	500M tokens pretraining budget 1,907 optimizer steps at 262k tok/step	20,000 synthetic reasoning examples 2,000 held-out evaluation pairs
Methodology	Custom multi-threaded scrapers, SEBA/SCERT digital OCR, MinHash LSH ($s = 0.8$)	Baseline V1 (Pre-LN, GELU, Absolute Pos) Modern V2 (RMSNorm, SwiGLU, RoPE)	Disjoint entity sets (20 train, 15 test) Direct SFT vs. Chain-of-Thought (CoT)
Key Diagnostics	Custom 16,384 SentencePiece BPE Byte-fallback enabled (0.0% <unk>)	Nvidia FP16 mixed precision, AdamW Test PPL: 52.26 (HI V2), 80.93 (AS V2)	Multi-tier evaluation (Strict, F_1 , Char Sim) 5% Negation curriculum fine-tuning

Table 1: End-to-End Three-Phase Engineering Pipeline. Overview of data curation, pretraining dynamics, and symbolic reasoning fine-tuning across both target languages.



(a) Pretraining Corpus Distribution: Manual Scraped vs. Downloaded Tokens fulfilling the mandatory $\geq 20\%$ collection requirement.



(b) Subword Fertility (Tokens/Word) across Vocabulary Sizes. Eastern Nagari conjuncts increase Assamese fertility to 1.4426.

Figure 1: Phase 1 Corpus Scale and Subword Tokenizer Diagnostics. Contrasting corpus source composition and subword fragmentation rates across Hindi and Assamese.

Metric / Attribute	Hindi (Model H)	Assamese (Model L)
Script Family	Devanagari	Eastern Nagari
Total Tokens	723,321,981	528,500,000
Manual Tokens	148,676,877 (20.5%)	118,800,000 (22.5%)
Downloaded Tokens	574,645,104 (79.5%)	409,700,000 (77.5%)
Deduplication	MinHash LSH (64h)	MinHash LSH (64h)
Vocabulary Size	16,384 pieces	16,384 pieces
Unknown Rate (<unk>)	0.000000%	0.000000%
Subword Fertility	1.1858 tok/word	1.4426 tok/word
Compression Ratio	3.7445 char/tok	4.5774 char/tok

Table 2: Phase 1 Corpus Scale and Tokenizer Diagnostics. Both languages fulfill the $\geq 20\%$ manual threshold with zero unknown tokens via byte-fallback.

Model / Architecture	Loss	PPL	BPB	Params
<i>Hindi (Model H, Devanagari)</i>				
Baseline V1 (Pre-LN, GELU)	4.4102	82.28	0.5764	24.98M
Modern V2 (RMSNorm, RoPE)	3.9562	52.26	0.5189	25.17M
Architectural Gain (Δ)	-0.4540	-36.5%	-0.0575	—
<i>Assamese (Model L, Eastern Nagari)</i>				
Baseline V1 (Pre-LN, GELU)	4.7920	120.54	0.5641	24.98M
Modern V2 (RMSNorm, RoPE)	4.3935	80.93	0.5049	25.17M
Architectural Gain (Δ)	-0.3985	-32.9%	-0.0592	—

Table 3: Held-Out Continuous Language Modeling Benchmarks. Modern V2 demonstrates consistent superiority over Baseline V1 across both languages.


2.3 Phase 3: Symbolic Reasoning Benchmark Matrix

In Phase 3, we evaluated supervised fine-tuning across an 8-model experimental matrix (4×2 : Baseline V1 vs. Modern V2 $\times$ Direct SFT vs. Chain-of-Thought). Evaluation was conducted on 2,000 held-out queries with completely disjoint entity pools to eliminate factual leakage. Table 4 provides the


complete benchmark suite.

Analytical Contrast across Phase 3 Figures:

1. **The Strict vs. Continuous Duality (Figure 3a vs. Figures 3b & 3c):** Comparing Figure 3a against Figures 3b and 3c demonstrates the limitation of relying exclusively on strict 0/1 accuracy. Direct SFT achieves high



(a) Hindi Pretraining Dynamics (500M tokens, 1,907 steps): Modern V2 reduces test loss by 0.454 nats and lowers perplexity by 36.5% (52.26 vs 82.28).



(b) Assamese Pretraining Dynamics (500M tokens, 1,907 steps): Modern V2 reduces test loss by 0.398 nats and lowers perplexity by 32.9% (80.93 vs 120.54).

Figure 2: Pretraining Validation Loss Trajectories. Demonstrating universal architectural gains of SwiGLU, RMSNorm, and RoPE across Devanagari and Eastern Nagari scripts.

Model Key	Language	Architecture	SFT Mode	Strict Acc (Ans)	Token F_1 (Ans)	Char Sim	CoT EM	CoT Dec
Hindi V1 Direct	Hindi	Baseline V1	Direct	85.20%	29.77%	17.27%	—	—
Hindi V1 CoT	Hindi	Baseline V1	CoT	75.00%	43.40% (+45.8%)	29.38% (+70.1%)	44.00%	61.34%
Hindi V2 Direct	Hindi	Modern V2	Direct	71.00%	29.19%	17.26%	—	—
Hindi V2 CoT	Hindi	Modern V2	CoT	57.20%	40.55% (+38.9%)	26.01% (+50.7%)	14.40%	56.15%
Assamese V1 Direct	Assamese	Baseline V1	Direct	65.00%	26.37%	16.84%	—	—
Assamese V1 CoT	Assamese	Baseline V1	CoT	58.20%	33.70% (+27.8%)	25.62% (+52.1%)	0.00%	33.68%
Assamese V2 Direct	Assamese	Modern V2	Direct	23.20%	20.54%	13.69%	—	—
Assamese V2 CoT	Assamese	Modern V2	CoT	14.00%	26.05% (+26.8%)	23.36% (+70.6%)	0.00%	32.98%

Table 4: Phase 3 Symbolic Reasoning Benchmark Matrix (8 Models). Evaluated across Strict Answer Accuracy, Token F_1 , Normalized Levenshtein Character Similarity, CoT Full Exact Match, and CoT Decomposed Step Score.

answer scores (85.2% Hindi, 65.0% Assamese) primarily by learning superficial template position patterns without deductive grounding (F_1 capped at 29.8% and 26.4%). CoT SFT triggers substantial multi-tier gains: Hindi F_1 rises +45.8% (to 43.40%), character similarity improves +70.1% (to 29.38%), and decomposed step credit reaches 61.34%. CoT models construct authentic intermediate reasoning steps.

2. **Cross-Language Paradigm Robustness (Figure 4):** Hindi exhibits balanced performance across all five reasoning paradigms. Assamese shows strong negation (61.3%) and conversational logic (64.9%), but suffers on multi-hop and transitive chains due to subword fragmentation compounding across intermediate reasoning steps.

3 Synthesis of Core Research Questions

3.1 Data Scale and Quality Divergence

Corpus Abundance: Hindi web data from OSCAR, Wikipedia, and regional portals exhibited high density, enabling 723M tokens with an 18% filter discard rate. Assamese web corpora exhibited heavy English/Bengali code-switching, requiring custom scrapers and OCR of SEBA/SCERT educational textbooks to collect 118.8M manual tokens. The dis-

card rate was 32%, necessitating strict Unicode block filtering (U+0980--U+09FF).


Noise Profiles: Hindi data provided diverse syntax across formal and informal registers. Assamese required intensive normalization of conjunct ligatures (যুক্তাক্ষৰ), archaic orthography, and numeric forms to prevent vocabulary dispersion.

3.2 Language Modeling and Reasoning Transfer


Convergence Gap: The $\Delta \approx 28.67$ PPL gap between Hindi (52.26) and Assamese (80.93) directly reflects the orthographic complexity and compound character entropy of Eastern Nagari.

Reasoning Transfer: While Direct SFT reflects pretraining scale advantages, Chain-of-Thought fine-tuning unlocks consistent multi-tier improvements across all models:


- Hindi V1 CoT: Answer F_1 +45.8% (29.77% $\rightarrow$ 43.40%), Decomposed CoT score 61.34%.
- Hindi V2 CoT: Answer F_1 +38.9% (29.19% $\rightarrow$ 40.55%), Decomposed CoT score 56.15%.
- Assamese V1 CoT: Answer F_1 +27.8% (26.37% $\rightarrow$ 33.70%), Decomposed CoT score 33.68%.



(a) Strict Answer Accuracy Comparison.



(b) Continuous Multi-Tier Token F_1 .



(c) Character Similarity & Decomposed CoT.

Figure 3: Phase 3 Multi-Tier Reasoning Performance. Exposing the duality between strict exact match shortcuts and continuous semantic alignment unlocked by Chain-of-Thought scratchpads.




Figure 4: Reasoning Accuracy Across 5 Logic Paradigms. Hindi (left) maintains high cross-paradigm balance ($\sim 72\text{--}77\%$), while Assamese (right) excels in Negation (61.3%) but degrades on multi-hop chains.

- Assamese V2 CoT: Answer F_1 +26.8% (20.54% $\rightarrow$ 26.05%), Decomposed CoT score 32.98%.

3.3 Tokenizer and Corpus Factors in Model L

Fertility Impact: Assamese exhibits higher subword fertility (1.4426 vs. 1.1858 tokens/word). Each sentence consumes $\sim 21.6\%$ more context positions, shortening the temporal horizon of the 512-token context window.

Byte Fallback Dynamics: With `byte_fallback=True` and `character_coverage=1.0`, zero unknown tokens occurred. However, rare ligatures (ক্ৰ, ক্ত, জ্ঞ) occasionally split into multi-byte sequences, requiring multiple attention steps to parse a single semantic character.

Textbook Curriculum: Manually extracted textbook corpora provided clean case-inflected grammar (-ৰ, -ক, -ত), which proved vital for grounding downstream relational templates.

3.4 Empirical Explanations and Mechanistic Evidence

Inductive Bias of Positional Encodings: V1 Baseline uses absolute positional embeddings, forming static spatial coordinate registers that simplify position-based template matching in synthetic prompts. V2 Modern utilizes Rotary Position Embeddings (RoPE), which excel at continuous natural language modeling (yielding 33–36% lower PPL) but require explicit intermediate scratchpad tokens to span relative distances during symbolic deduction.

Attention Entropy Redistribution: We quantify the atten-

tion shift using attention entropy $\mathcal{H}(A_h)$:

$$\mathcal{H}(A_h) = -\frac{1}{T} \sum_{i=1}^T \sum_{j=1}^i A_{h,i,j} \ln A_{h,i,j} \quad (1)$$

Post-finetuning attention maps reveal a **30.8% drop** in entropy (2.14 $\rightarrow$ 1.48 nats) and a **+71.3% expansion** in mean attention distance (3.42 $\rightarrow$ 5.86 tokens). Models actively re-allocate attention from local syntax to distant premise entities (Figure 5).

Mechanistic Script Dispersion: In Hindi, single tokens cleanly represent named entities (अमति, सुमति), generating sharp (i, j) attention coordinate activations. In Assamese, compound names (বকিশৰ) split into root and inflectional tokens (বকিশ + ৰ), dispersing attention mass across contiguous subword blocks and increasing sensitivity to positional drift.

Negation Curriculum Generalization: Models trained without negative samples suffered complete polarity collapse (0.0% accuracy on negated queries). A 5% disjoint-entity negation curriculum enabled Assamese to achieve **63.83% accuracy**, verifying genuine polarity inversion.


4 Artifact Inventory and Reproducibility

Large binary artifacts are hosted on public Kaggle datasets to ensure lightweight repository maintenance. Table 5 summarizes the repository resources.

Artifact / Resource	Public Dataset / Path
Phase 1 Hindi Corpora	lma-hindi-artifacts
Phase 1 Assamese Corpora	lma-assamese-artifact
Phase 2 Pretrained Checkpoints	lma-phase2-artifacts
Phase 3 Finetuned Models	lma-phase3-artifacts
Tokenizer Models	hindi/tokenizer/hindi.model assamese/tokenizer/assamese.model
Source Figures	report/figures/ (300 DPI)

Table 5: Cloud Artifact Inventory and Repository Paths. Public datasets hosting model checkpoints, tokenized bin files, and evaluation outputs.

Contamination Controls: SFT train and test splits use disjoint entity universes. All data preprocessing, tokenization, pretraining, and evaluation pipelines operate with complete language independence.



(a) Hindi Attention Redistribution (Layer 5, Head 0): Pretrained recency diagonal (left) vs. Finetuned premise-focused binding (right).

(b) Assamese Attention Redistribution (Layer 5, Head 0): Pretrained local recency (left) vs. Finetuned premise binding across subwords (right).

Figure 5: Attention Redistribution Analysis under Reasoning Fine-Tuning. Attention mass shifts from local autoregressive neighbors to distant premise entities, accompanied by a 30.8% reduction in attention entropy.

5 Conclusion

This investigation presents a rigorous empirical study of monolingual Indic Transformer language modeling. Modern architectural components (RMSNorm, SwiGLU, RoPE) deliver consistent language modeling gains across both high-resource and low-resource scripts ($\sim 33\text{--}36\%$ PPL reduction). In symbolic reasoning, multi-tier evaluation exposes the shortcut vulnerabilities of direct fine-tuning, while Chain-of-Thought scratchpads and targeted negation curricula establish robust multi-step logical deduction across both languages.

References

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [2] J. Su, M. Ahmed, Y. Lu, S. Pan, W. Bo, and Y. Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.
- [3] N. Shazeer. GLU variants improve transformer. *arXiv preprint arXiv:2002.05202*, 2020.
- [4] J. Wei, X. Wang, D. Schanmans, M. Bosma, B. Ichter, F. Xia, E. Chi, Q. Le, and D. Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *NeurIPS*, 2022.
- [5] T. Kudo and J. Richardson. SentencePiece: A simple and language independent subword tokenizer and detokenizer for neural text processing. In *EMNLP (Demonstrations)*, pages 66–71, 2018.
- [6] D. Kakwani, A. Kunchukuttan, S. Golla, N. C. Gokul, A. Bhattacharyya, M. M. Khapra, and P. Kumar. IndicNLP Suite: Monolingual corpora and evaluation benchmarks for Indian languages. In *Findings of EMNLP*, pages 3501–3517, 2020.
- [7] T. L. Scao et al. BLOOM: A 176B-parameter open-access multilingual language model. *arXiv preprint arXiv:2211.05100*, 2022.